


# Leveraging the Distributed Model on Demand Platform to Work With Community Models in the Next Generation Water Resources Modeling Framework

Robert Bartel, Austin Raney, Christopher Tubbs, Nels J. Frazier, Trey Flowers, Shengting Cui

NWS/Office of Water Prediction/National Water Center; Lynker Technologies, LLC

## 1 Intro - What is DMOD?



- Prototype platform for specialized, scalable compute environments (especially with ngen)
- **Key Feature:** abstract and manage compute infrastructure and environments


OPEN

## 2 What's Needed for ngen



OPEN

## 3 Using Docker Containers




OPEN

## 4 Missing BMI Modules



OPEN

## 5 Guided Customization




Customization Directory

- Special Git-ignored customization

OPEN

## 6 An Example: A Forked LSTM BMI Module



OPEN

## 7 Summary and Future

### What DMOD Does

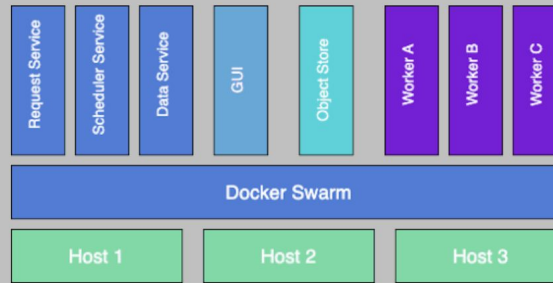
- Automates performance of typical steps for creating consistent compute environments for ngen via Docker images
- Enables local customization to incorporate external BMI modules into image builds without modifying distributed source code

### Future Ideas

- Still room for improvements
- Both in and beyond DMOD

OPEN

# 1 Intro - What is DMOD?



- Prototype software suite to run specialized, scalable compute environments

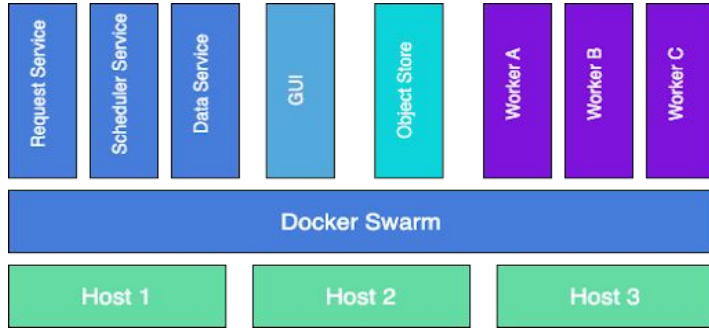
- **Key Feature: abstract and manage compute infrastructure**

Distributed Model on Demand (DMOD) is an extensible suite of software tools for creating and running specialized compute environments. The primary goal for DMOD is to make it easier to develop, test, and experiment with scientific models, with particular emphasis on models run through *ngen*, the model engine of the NextGen framework.

More information on DMOD, as well as the source code, can be found here:

<https://github.com/NOAA-OWP/DMOD>

# 1 Intro - What is DMOD?



- Prototype software suite to run specialized, scalable compute environments
- **Key Feature: abstract and manage compute infrastructure**

Distributed Model on Demand (DMOD) is an extensible suite of software tools for creating and running specialized compute environments. The primary goal for DMOD is to make it easier to develop, test, and experiment with scientific models, with particular emphasis on models run through ngen, the model engine of the NextGen framework.

More information on DMOD, as well as the source code, can be found here:

<https://github.com/NOAA-OWP/DMOD>

## 2 What's Needed for ngen

### Language Support



### Dependencies



### Software Tools

## 2 What's Needed for ngen



### Software Tools



There are a number of things needed for running ngen within a compute environment:

- In order to compile and run ngen and BMI modules, language support is needed for several different programming languages.
- Software development tools are required to get and build ngen, dependency and BMI module source code - e.g., Git, GCC, CMake.
- Before ngen can be compiled, several other dependencies must be installed, from packages or by compiling source code; e.g., netCDF, MPICH, and SQLite.
- ngen itself must be compiled.
- BMI modules must be compiled and/or installed, depending on the programming language.

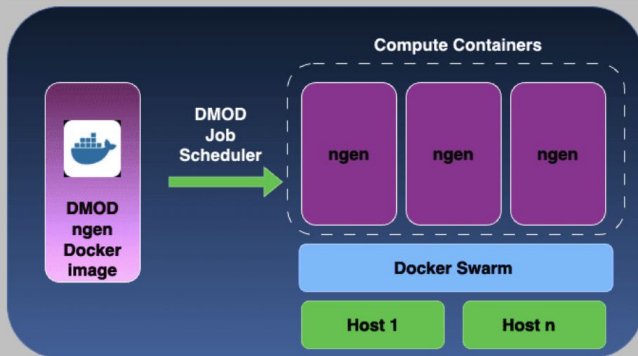

## 2 What's Needed for ngen



There are a number of things needed for running ngen within a compute environment:

- In order to compile and run ngen and BMI modules, language support is needed for several different programming languages.
- Software development tools are required to get and build ngen, dependency and BMI module source code - e.g., Git, GCC, CMake.
- Before ngen can be compile, several other dependencies must be installed, from packages or by compiling source code; e.g., netCDF, MPICH, and SQLite.
- ngen itself must be compiled.
- BMI modules must be compiled and/or installed, depending on the programming language.

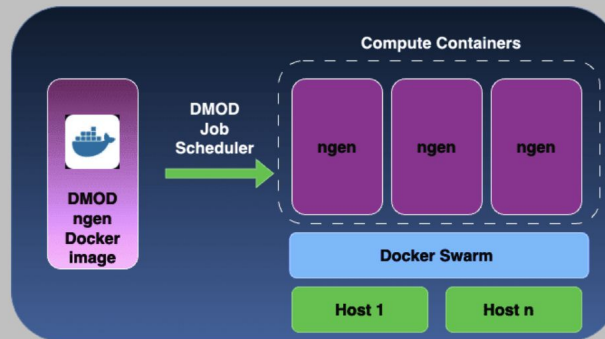
### 3 Using Docker Containers



#### DMOD's Approach

- Maintains compute environment infrastructure as source code via *Dockerfiles*
- Includes custom tools to locally build Docker *images* from these Dockerfiles
- Has services that execute ngen by starting Docker *containers* from those images

### 3 Using Docker Containers



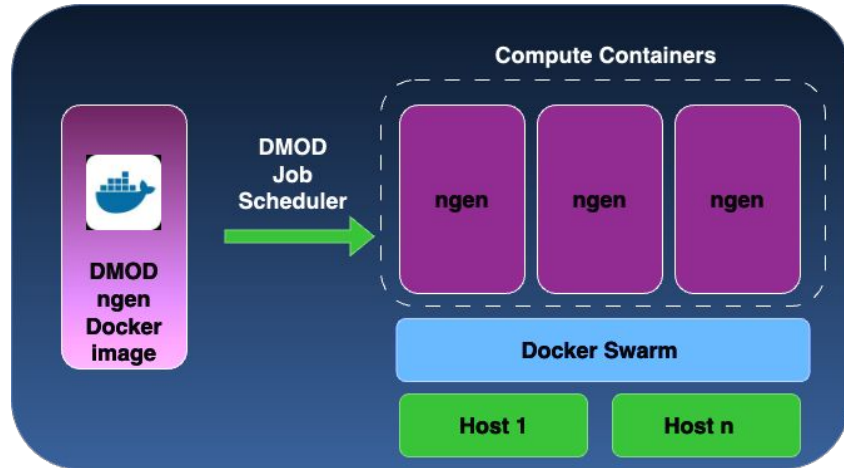

#### DMOD's Approach

- Maintains compute environment infrastructure as source code via *Dockerfiles*
- Includes custom tools to locally build Docker *images* from these Dockerfiles
- Has services that execute ngen by starting Docker *containers* from those images

#### Key Benefits

- Users don't need to figure out how to compile ngen
- Users get consistent, re-creatable compute environments

# 3 Docker Containers



## DMOD's Approach

- Maintains compute environment infrastructure as source code via Dockerfiles
- Includes custom tools to locally build Docker images from these Dockerfiles
- Has services that execute ngen by starting Docker containers from those images

## Key Benefits

- Users don't need to figure out how to compile ngen
- Users get consistent, re-creatable compute environments



## 4 Missing BMI Modules



### The Problem

BMI prevents DMOD from hard-coding integration of all modules of interest into Dockerfile source code.

### Technical Limits

- Tremendous advantage: ngen can use essentially any external model adapted to use the BMI
- Consequence: simply cannot know about every BMI module everywhere
- Even best-effort would inevitably lead to **conflicts** with **dependencies** while still never being enough for every situation

### Another Consideration



## 4 Missing BMI Modules



### The Problem

BMI prevents DMOD from hard-coding integration of all modules of interest into Dockerfile source code.

### Technical Limits

- Tremendous advantage: ngen can use essentially any external model adapted to use the BMI
- Consequence: simply cannot know about every BMI module everywhere
- Even best-effort would inevitably lead to **conflicts** with **dependencies** while still never being enough for every situation

### Another Consideration



Further, there could be licensing and rights restrictions that prevent bundling of certain software without changing DMOD's own licensing.



# 4 Limits on Unlimited



## The Problem

BMI prevents DMOD from hard-coding integration of all modules of interest into Dockerfile source code.

## Technical Limits


- Tremendous advantage: ngen can use essentially any external model adapted to use the BMI
- Consequence: simply cannot know about every BMI module everywhere
- Even best-effort would inevitably lead to conflicts with dependencies while still never being enough for every situation

## Another Consideration



Further, there could be licensing and rights restrictions that prevent bundling of certain software without changing DMOD's own licensing.

## 5 Guided Customization



### Customization Directory

- Special Git-ignored customization directory: [docker/main/ngen/customize/](https://github.com/docker/main/ngen/customize/)
- A "bucket" in which to put necessary file(s) for supported customization method(s)
- Can also place supplementary files here; entire contents copied into image build working directory

### Methods for Customization

- Three ways to inject customization (documented at [docker/main/ngen/customize/README.md](https://github.com/docker/main/ngen/customize/README.md)):
  - A pip requirements.txt file
  - A list of repos to auto-build with CMake
  - A Bash script for fine-grained control

### What Happens

Directory contents only locally present

## 5 Guided Customization



### Customization Directory

- Special Git-ignored customization directory: [docker/main/ngen/customize/](https://github.com/docker/main/ngen/customize/)
- A "bucket" in which to put necessary file(s) for supported customization method(s)
- Can also place supplementary files here; entire contents copied into image build working directory


### Methods for Customization

- Three ways to inject customization (documented at [docker/main/ngen/customize/README.md](https://github.com/docker/main/ngen/customize/README.md)):
  - A pip requirements.txt file
  - A list of repos to auto-build with CMake
  - A Bash script for fine-grained control

### What Happens

- Directory contents *only locally present*
  - *Not committed to the repository*
- When present, these get used automatically by DMOD image build tools
- They trigger execution of the customization steps at a specific stage of the ngen Docker image build

# 5 Enable Customization



## Customization Directory

- Special Git-ignored customization directory: `docker/main/ngen/customize/`
- A "bucket" in which to put necessary file(s) for supported customization method(s)
- Can also place supplementary files here; entire contents copied into image build working directory


## Methods for Customization

- Three ways to inject customization (documented at `docker/main/ngen/customize/README.md`):
  - A `pip requirements.txt` file
  - A list of repos to auto-build with CMake
  - A Bash script for fine-grained control

## What Happens

- Directory contents only locally present
  - Not committed to the repository
- When present, these get used automatically by DMOD image build tools
- They trigger execution of the customization steps at a specific stage of the ngen Docker image build

## 6 An Example: A Forked LSTM BMI Module



To illustrate, consider an LSTM BMI module.

The source code can be found at [https://github.com/robertbartel/owp\\_lstm](https://github.com/robertbartel/owp_lstm)

*Important: not pre-integrated into Docker image, not (directly) from OWP's public repo.*

### The Steps

Create a local `docker/main/ngen/customize/customize.sh` file:

```
1 #!/bin/bash
2
3 # Activate the Python virtual environment
4 source /dmod/venv/bin/activate
5
6 # Install the required dependencies for the module, via OS manager and Pip
7 dnf install -y libgomp libomp libomp-devel
8 pip install --no-cache-dir torch==2.3.1+cpu -f https://download.pytorch.org/whl/torch_stable.html
9 pip install --no-cache-dir wheel bmpi bokeh jupyter matplotlib netcdf4 pandas ruamel.yaml
10 pip install --no-cache-dir -U --force-reinstall xarray==0.16.0
```

## 6 An Example: A Forked LSTM BMI Module

```
1 #!/bin/bash
2
3 # Activate the Python virtual environment
4 source /dmod/venv/bin/activate
5
6 # Install the required dependencies for the module, via OS manager and Pip
7 dnf install -y libgomp libomp libomp-devel
8 pip install --no-cache-dir torch==2.3.1+cpu -f https://download.pytorch.org/whl/torch_stable.html
9 pip install --no-cache-dir wheel bmipy bokeh jupyter matplotlib netcdf4 pandas ruamel.yaml
10 pip install --no-cache-dir -U --force-reinstall xarray==0.16.0
11
12 # Download (with Git) and installed the LSTM BMI module
13 git clone https://github.com/robertbartel/owp_lstm.git /dmod/lstm
14 pip install --no-deps /dmod/lstm
15
16 # Setup the pre-trained model file and the training config
17 cp -a /dmod/lstm/trained_neuralhydrology_models/hourly_slope_mean_precip_temp /dmod/bmi_module_data/lstm
18 sed -i.bak 's/\.\/trained_neuralhydrology_models\/hourly_slope_mean_precip_temp\/\dmod\/bmi_module_data\/lstm/' \
19 /dmod/bmi_module_data/lstm/config.yml
20 chown -R mpi:mpi /dmod/bmi_module_data/lstm
21
22 # Do some cleanup to minimize Docker image size
23 dnf clean -y all
24 deactivate
25 rm -rf /dmod/lstm
26
~
~
~
```


19,4 All

Then rebuild the Docker image using standard DMOD helper script:

```
./scripts/control_stack.sh --build-args "ngen" main build push
```

With this, we can run LSTM prediction for 1 month on about 18,000 catchments (VPU01) in about 16 minutes. This was on a PC with 13th Gen Intel i7 processor, using 16 of its cores.

# 6 An Example: A Forked LSTM BMI Module



To illustrate, consider an LSTM BMI module.

The source code can be found at [https://github.com/robertbartel/owp\\_lstm](https://github.com/robertbartel/owp_lstm)

*Important: not pre-integrated into Docker image, not (directly) from OWP's public repo.*

## The Steps

Create a local `docker/main/ngen/customize/customize.sh` file:

```
1 #!/bin/bash
2
3 # Activate the Python virtual environment
4 source /dmod/venv/bin/activate
5
6 # Install the required dependencies for the module, via OS manager and Pip
7 /dmod/install -y libgomp libomp libomp-devel
8 pip install --no-cache-dir setuptools==5.7.2 -- https://download.pytorch.org/whl/torch_stable.html
9 pip install --no-cache-dir wheel==0.37.0 -- https://pypi.org/project/wheel/
10 pip install --no-cache-dir numpy==1.21.0 -- https://pypi.org/project/numpy/
11 pip install --no-cache-dir pandas==1.3.0 -- https://pypi.org/project/pandas/
12
13 # Download (with Git) and installed the LSTM BMI module
14 git clone https://github.com/robertbartel/owp_lstm.git /dmod/lstm
15 pip install --no-deps /dmod/lstm
16
17 # Setup the pre-trained model file and the training config
18 cp -a /dmod/lstm/trained_neuralhydrology_models/hourly_slope_mean_precip_temp /dmod/lstm_module_data/lstm
19 sed -i.bak -e 's|/dmod/lstm/trained_neuralhydrology_models/hourly_slope_mean_precip_temp|/dmod/lstm_module_data/lstm|' \
20     /dmod/lstm_module_data/lstm/config.yml
21
22 # Do some cleanup to minimize Docker image size
23 def clean -y all
24 deactivate
25 rm -rf /dmod/lstm
```

Then rebuild the Docker image using standard DMOD helper script:

```
./scripts/control_stack.sh
--build-args "ngen" main build push
```

With this, we can run LSTM prediction for 1 month on about 18,000 catchments (VPU01) in about 16 minutes. This was on a PC with 13th Gen Intel i7 processor, using 16 of its cores.

# 7 Summary and Future

## What DMOD Does

- Automates performance of typical steps for creating consistent compute environments for ngen via Docker images
- Enables local customization to incorporate external BMI modules into image builds without modifying distributed source code

## Future Ideas

- Still room for improvements
  - Both in and beyond DMOD
  - More interoperable ngen+BMI ecosystem
- This makes it easier for a model *you* know
  - Still can be tricky for one “off the shelf”
  - \* Tried to incorporate the [SUMMA](#) module
  - Could not in time constraints
    - Troubles independent of DMOD
- Integratration with the tools for BMI init config generation would be beneficial
- Need more tools and standards/conventions to “package” models more effectively and consistently
  - Standardizing versions and dependencies
  - Standardizing builds and installations
  - Standardizing paths
- Consider separating development of ngen Docker images into repo focused on them



# 7 Summary and Future

## What DMOD Does

- Automates performance of typical steps for creating consistent compute environments for ngen via Docker images
- Enables local customization to incorporate external BMI modules into image builds without modifying distributed source code

## Future Ideas

- Still room for improvements
  - Both in and beyond DMOD
  - More interoperable ngen+BMI ecosystem
- This makes it easier for a model you know
  - Still can be tricky for one “off the shelf”
  - Tried to incorporate the SUMMA module
  - Could not in time constraints
    - Troubles independent of DMOD
- Integration with the tools for BMI init config generation would be beneficial
- Need more tools and standards/conventions to “package” models more effectively and consistently
  - Standardizing versions and dependencies
  - Standardizing builds and installations
  - Standardizing paths
- Consider separating development of ngen Docker images into repo focused on them