



OWP | OFFICE OF
WATER
PREDICTION

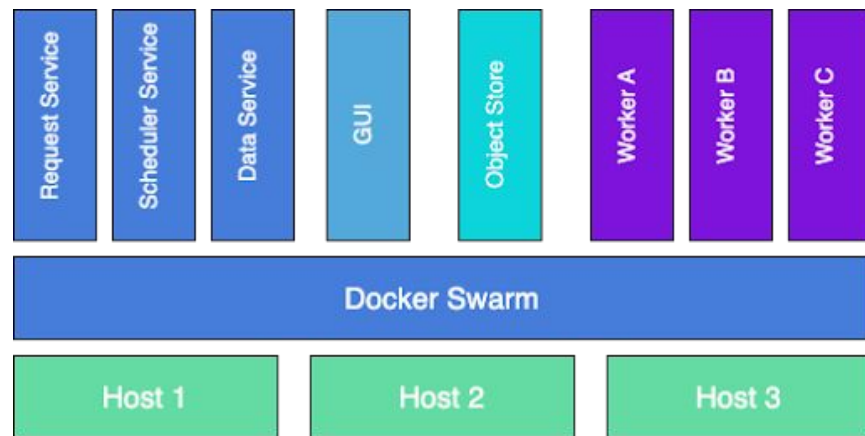
Leveraging the Distributed Model on Demand Platform to Work With Community Models in the Next Generation Water Resources Modeling Framework

*Robert Bartel¹, Austin Raney¹, Christopher Tubbs¹, Nels J. Frazier¹, Trey Flowers²,
Shengting Cui¹*



1. Lynker Technologies, LLC; NOAA Affiliate
2. NWS/Office of Water Prediction/National Water Center

Background: What is DMOD?



- Prototype platform for facilitating running ngen
- **Key Feature: abstract and manage compute infrastructure and environments**

Full details can be found at:

<https://github.com/NOAA-OWP/DMOD>



OWP | OFFICE OF
WATER
PREDICTION

What is in an ngen Compute Environment?



What is in an ngen Compute Environment?

Language Support



What is in an ngen Compute Environment?

Language Support



Dependencies

What is in an ngen Compute Environment?

Language Support



Software Tools and Builds



Dependencies

What is in an ngen Compute Environment?

Language Support



DMOD was designed to handle some of this burden for users.

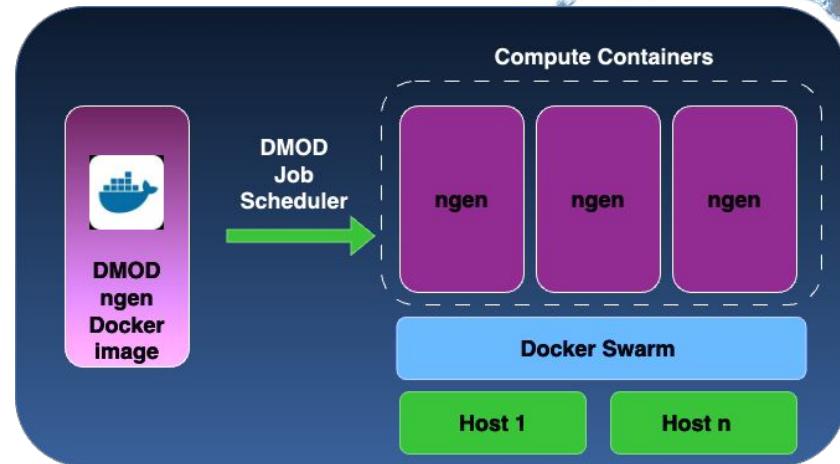


Software Tools and Builds



Dependencies

Package Environments in Docker Containers



- Compute infrastructure as source code
- Tools to *locally* build images and run ngen in containers
- No figuring out how to compile ngen from scratch
- Consistent, reproducible environments

What About BMI Modules?

- Similar dependency and install needs as ngen
- Image source code bundled with some OWP-developed BMI modules
- Works well for many use cases
- Eventually more will be needed

What About BMI Modules?

- Similar dependency and install needs as ngen
- Image source code bundled with some OWP-developed BMI modules
- Works well for many use cases



What About BMI Modules?

- Similar dependency and install needs as ngen
- Image source code bundled with some OWP-developed BMI modules
- Works well for many use cases
- Eventually more modules will be needed
- ngen designed to work with community-developed modules
- DMOD needs to provide the same flexibility



What About BMI Modules?

- Image source code bundled with some OWP-developed BMI modules
- Works well for many use cases
- Eventually more will be needed
- Why?



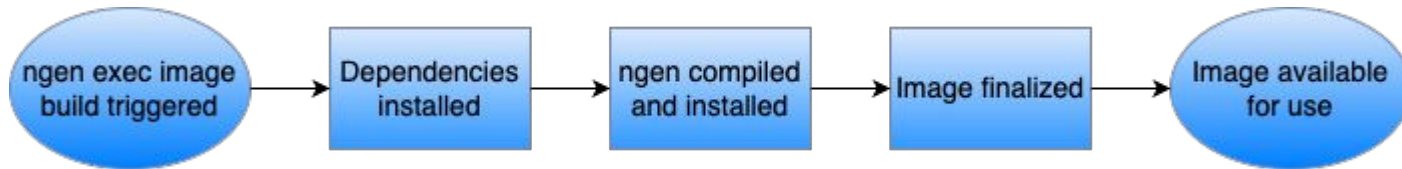
The Solution: Guided Customization

- Not practical to bundle every possible module
- Instead give users ways to customize Docker image
- Lets users insert BMI modules **they** need
- Three different guided methods supported
- Utilizes a special directory in local DMOD file structure
 - `docker/main/ngen/customize`
 - A “bucket” for customization files
 - Ignored by Git
 - Includes a README.md file with documentation

Providing Guided Customization

- Not practical to bundle every possible module
- Instead give users ways to customize Docker image
- Lets users insert BMI modules **they** need

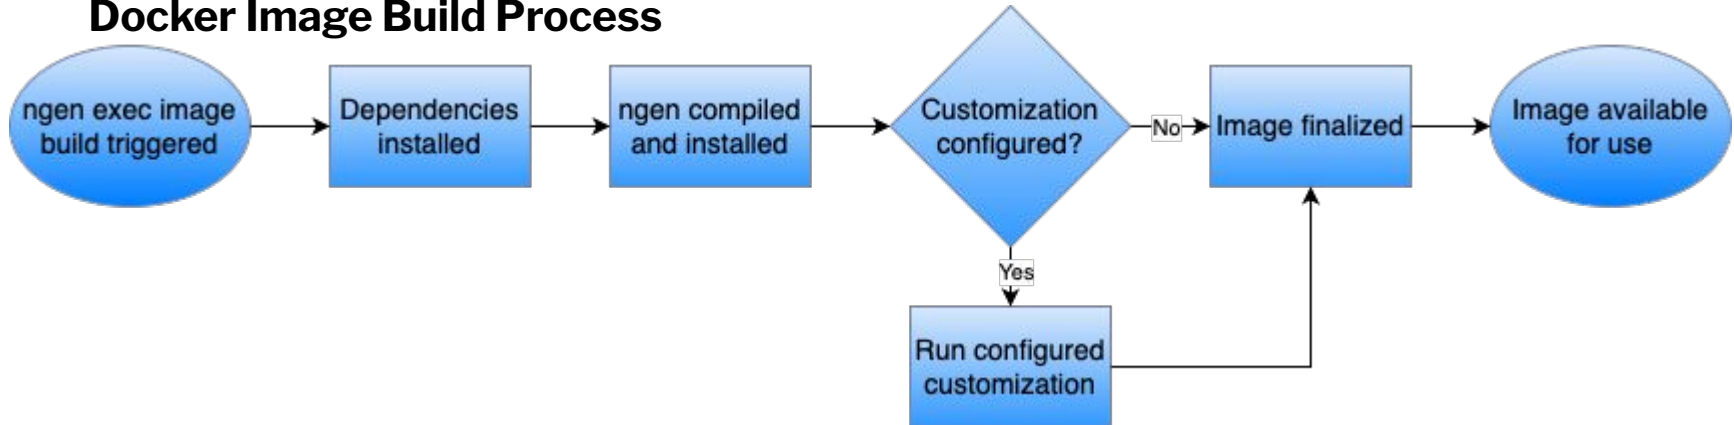
Docker Image Build Process



Providing Guided Customization

- Not practical to bundle every possible module
- Instead give users ways to customize Docker image
- Lets users insert BMI modules **they** need

Docker Image Build Process



Configuring Customization

Available Methods

- A pip requirements file of Python packages to install
- A list of Git repos to auto-build with CMake
- A Bash script for fine-grained control



Configuring Customization



Customize Directory

- Files placed in special directory in local DMOD file structure
 - `docker/main/ngen/customize`
- Ignored by Git
 - Only exists if locally added by user

How it Works

- User adds customization file(s) to directory in local copy of DMOD
- When present, contents used automatically by DMOD image build tools
 - No other specialized action needed
 - Presence triggers execution of customization steps at specific build stage

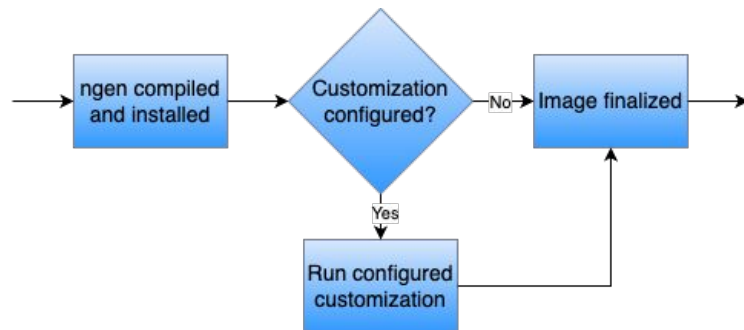
Configuring Customization

Customize Directory

- Files placed in special directory in local DMOD file structure
 - `docker/main/ngen/customize`
- Ignored by Git
 - Only exists if locally added by user

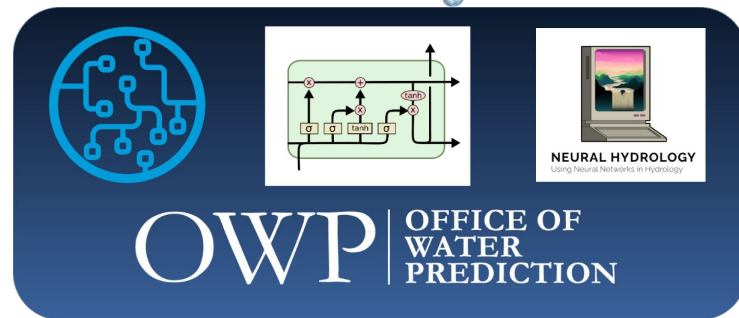
How it Works

- User adds customization file(s) to directory in local copy of DMOD
- When present, contents used automatically by DMOD image build tools
 - No other specialized action needed
 - Presence triggers execution of customization steps at specific build stage



An Example: A Fork of LSTM

- As an example, consider an LSTM BMI module
- Source code available here:
 - https://github.com/robertbartel/owp_lstm
- Analogous to arbitrary community BMI module
 - Not directly from OWP public repo
 - Not pre-integrated into image
- Because some special handling of dependencies is needed, we will use the Bash script



The Steps

Step 1

Create `docker/main/ngen/customize/customize.sh`

Step 2

(Re)build the Docker image using DMOD's *control_stack.sh* helper script:

```
./scripts/control_stack.sh --build-args "ngen" main build push
```

Step 3

Run a DMOD ngen job as normal, using the module in formulations

An Example customize.sh File

```
1 #!/bin/bash
2
3 # Activate the Python virtual environment
4 source /dmod/venv/bin/activate
5
6 # Install the required dependencies for the module, via OS manager and Pip
7 dnf install -y libgomp libomp libomp-devel
8 pip install --no-cache-dir torch==2.3.1+cpu -f https://download.pytorch.org/whl/torch_stable.html
9 pip install --no-cache-dir wheel bmipy bokeh jupyter matplotlib netcdf4 pandas ruamel.yaml
10 pip install --no-cache-dir -U --force-reinstall xarray==0.16.0
11
12 # Download (with Git) and installed the LSTM BMI module
13 git clone https://github.com/robertbartel/owp_lstm.git /dmod/lstm
14 pip install --no-deps /dmod/lstm
15
16 # Setup the pre-trained model file and the training config
17 cp -a /dmod/lstm/trained_neuralhydrology_models/hourly_slope_mean_precip_temp /dmod/bmi_module_data/lstm
18 sed -i.bak 's/\.\/trained_neuralhydrology_models\/hourly_slope_mean_precip_temp\/\dmod\/bmi_module_data\/lstm/' \
19     /dmod/bmi_module_data/lstm/config.yml
20 chown -R mpi:mpi /dmod/bmi_module_data/lstm
21
22 # Do some cleanup to minimize Docker image size
23 dnf clean -y all
24 deactivate
25 rm -rf /dmod/lstm
26
~
~
~
```

19,4 All

The Steps (continued)

Step 2

(Re)build the Docker image using DMOD's *control_stack.sh* helper script:

```
./scripts/control_stack.sh --build-args "ngen" main build push
```

Step 3

Run a DMOD ngen job as normal, using LSTM in formulations.

High Level Results

- Ran a simple ngen job through DMOD using LSTM
- Trained model params:
 - Slope, Elevation, Precipitation, Temperature
- 1 month
 - January 2016
- Hydrofabric VPU-01
 - ~18,000 catchments
- 16 CPU cores
 - Intel 13th Gen i7
- Available memory: ~90 GB memory
 - Max used was ~80 GB
- Job time: ~16 minutes

Summary and Future

What DMOD Does

- Automates steps to create consistent compute environments via Docker images and containers
- Enables guided local customization to integrate external BMI modules into images without modifying distributed source code

Future Ideas

- More tools and standards/conventions to “package” models
- Separate image building from DMOD and unify with other work



OWP | OFFICE OF
WATER
PREDICTION



Thank You!



For more information:

Robert “Bobby” Bartel



robert.bartel@noaa.gov



<https://water.noaa.gov>